

Praktikum Rechnernetze

Versuch 1: Troubleshooting TCP/IP

Elia Wüstner, Daniel Hiller, Felix Pojtinger, Jakob Waibel

19.10.2021

Introduction

Contributing

These study materials are heavily based on professor Kiefer's "Praktikum Rechnernetze" lecture at HdM Stuttgart.

Found an error or have a suggestion? Please open an issue on GitHub (github.com/pojntfx/uni-netpractice-notes):



Figure 1: QR code to source repository

If you like the study materials, a GitHub star is always appreciated :)

License



Figure 2: AGPL-3.0 license badge

Uni Network Practice Notes (c) 2021 Felix Pojtinger and contributors

SPDX-License-Identifier: AGPL-3.0

IP-Subnetz-Berechnung

Ergänzen Sie die Tabelle

IP-Adresse	SN-Mask	Klasse	Netz- adresse	Anzahl Subnetze	Broadcast- Adresse	Anzahl Hosts	Vorheriges Netz	nachgelag. Netz
14.21.4.210	255.255.128.0	A	14.21.0.0	512	14.21.127.255	32.768	14.20.0.0	14.21.128.0
184.16.12.80	255.255.255.224	B	184.16.12.64	2048	184.16.12.95	30	184.16.12.32	184.16.12.96
143.62.67.32	255.255.255.240	B	143.62.67.32	4096	143.62.67.47	16	143.62.67.16	143.62.67.80
264.12.14.81	255.255.192.0	/	/	/	/	/	/	/
192.168.1.42	255.255.255.0	C	192.168.1.0	1	192.168.1.255	254	/	/
10.15.119.237	255.255.255.252	A	10.15.119.232	4096	10.15.119.239	2	10.15.119.232	10.15.119.240

184.16.12.80 → Class B

255.255.255.224

$8 = 2^3 = 8 \rightarrow 3 \rightarrow 127 \rightarrow 184.16.12.127$ 11000

255.255.255.11110000 → 224

184.16.12.011010000 → 80

011000000 → 64 → 184.16.12.64 | Network address

011011111 → 95 → 184.16.12.95 | Broadcast address

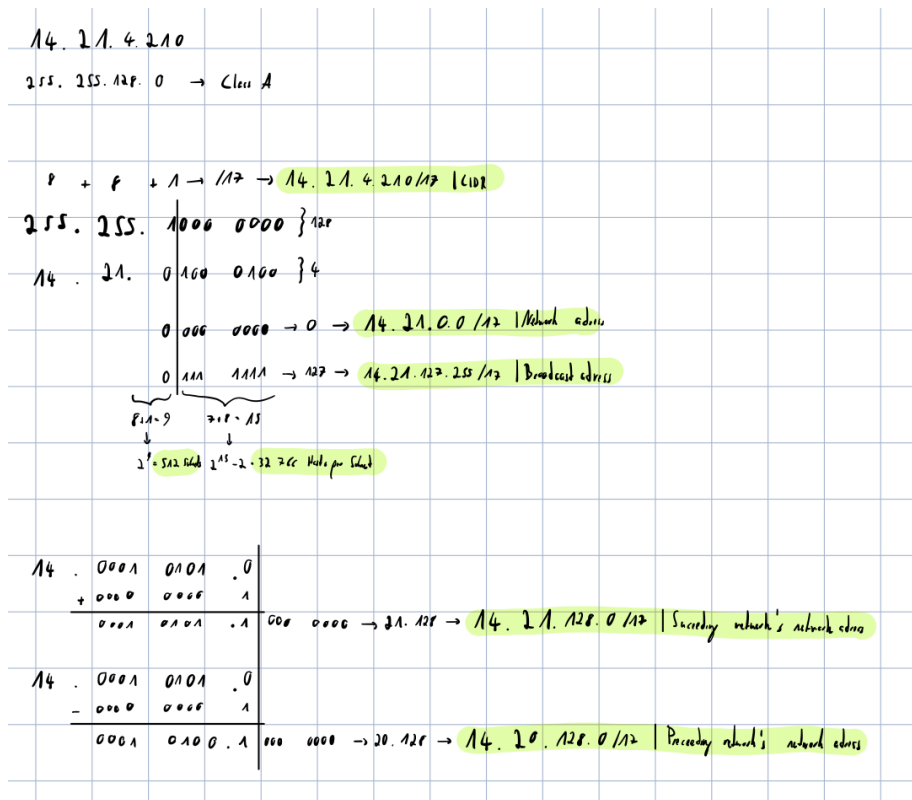
$2^3 = 8$
 $2^4 = 16$
 $2^5 = 32$
 $2^6 = 64$
 $2^7 = 128$
 $2^8 = 256$
 $2^9 = 512$
 $2^{10} = 1024$
 $2^{11} = 2048$
 $2^{12} = 4096$
 $2^{13} = 8192$
 $2^{14} = 16384$
 $2^{15} = 32768$
 $2^{16} = 65536$
 $2^{17} = 131072$
 $2^{18} = 262144$
 $2^{19} = 524288$
 $2^{20} = 1048576$
 $2^{21} = 2097152$
 $2^{22} = 4194304$
 $2^{23} = 8388608$
 $2^{24} = 16777216$
 $2^{25} = 33554432$
 $2^{26} = 67108864$
 $2^{27} = 134217728$
 $2^{28} = 268435456$
 $2^{29} = 536870912$
 $2^{30} = 1073741824$
 $2^{31} = 2147483648$

255.255.255.11110000 → 96 → 184.16.12.96 | Network address

255.255.255.11110000 → 96 → 184.16.12.96 | Network address

143.62.67.32

255.255.255.240 → Class B



Tools des OS

IP-Konfiguration

Überprüfen Sie zunächst die Netzkonfiguration Ihres PC. IP-Adresse, Subnetzmaske, Default-Gateway und DNS-Server Erfragen Sie den Klartextnamen Ihres PC.

IP-Adresse: 142.62.66.5

Subnetzmaske: 255.255.255.0

Default-Gateway: 141.62.66.250

DNS-Server: 141.62.66.250

Klartextnamen: rm05

Wie können Sie die korrekte Installation der Netzwerkkarten-Treiber testen?

```
$ lspci
2 # ...
00:1f.6 Ethernet controller: Intel Corporation Ethernet
      Connection (2) I219-LM
```

```

4 # ...
  $ find /sys | grep drivers.*00:1f.6
6 # ...
  /sys/bus/pci/drivers/e1000e/0000:00:1f.6

```

Testen Sie die DNS-Namensauflösung mit nslookup

Wir verwenden an dieser Stelle dig, da nslookup deprecated ist. Die Option +noall entfernt alle Display-Flags und +answer zeigt dann nur die Antwortsektion des Outputs an.

```

$ dig +noall +answer +multiline www.hdm-stuttgart.de
2 www.hdm-stuttgart.de. 3553 IN A 141.62.1.53
  www.hdm-stuttgart.de. 3553 IN A 141.62.1.59

```

Wir erhalten zwei Ergebnisse auf unsere Anfrage. Das könnte daran liegen, dass die HdM zur Lastenaufteilung zwei Webserver einsetzt.

Anschluss des PC an das Labornetz

Betrachten Sie die Verbindungen der Labor-Switches untereinander. Welche Wege können Sie erkennen?

Folgende Verbindungen konnten erkannt werden:




Figure 3: Unser Computer ist an die RJ-45-Buchse 1-01 angeschlossen. Das Kabel der Buchse führt dann in den Netzwerkschrank

Wenn die Verbindung am Patch-Panel zu 1-01 unterbrochen wird, so verliert die Netzwerkkarte die Verbindung, was der Kernel-Buffer bestätigt:




Figure 4: Auf diesem Bild ist der Netzwerkschrank zu sehen. Man sieht hier das Patchfeld, an welchem die 1-01 angeschlossen ist. Vom Patchfeld führt ein weiteres LAN-Kabel (CAT-5e) zu einem Switch.



Figure 5: Der Switch ist dann mit dem hier zu sehenden Router verbunden. Der Router führt dann zur restlichen Infrastruktur des Hauses bzw. zum Internet.

```

1 $ dmesg -w
# ...
3 [ 6.048643] e1000e 0000:00:1f.6 enp0s31f6: NIC Link is Up
    1000 Mbps Full Duplex, Flow Control: None
  [ 1360.221984] e1000e 0000:00:1f.6 enp0s31f6: NIC Link is Down
5 # ...

```

Verfolgen Sie den im Netzwerkschrank gepatchten Weg, auf dem die Pakete Ihres Rechners zum Router gelangen

Wie schon an den Bildern vorher illustriert lässt sich folgender Weg ableiten:

```

1 Patch-Feld -> Switch -> Router -> Rest der Infrastruktur

```

Verfolgen Sie den Weg, auf dem die Pakete Ihres Rechners den gegenüberliegenden Netzwerkschrank erreichen

Warum ist im Netzwerkschrank wohl ein Hub installiert?

Es ist ein Hub installiert, sodass die verschiedenen Nodes im LAN-Netzwerk miteinander kommunizieren können. Dies ermöglicht zudem auch einfacheres Debugging über Sniffing.

Überprüfung der korrekten Installation

Sehen Sie sich die IP-Konfiguration Ihres Rechners an durch Eingabe von ipconfig bzw. ipconfig/all in der DOS-Box.

ipconfig ist deprecated, es wird stattdessen ip verwendet.

```

$ ip a
2 1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state
    UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
4     inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
6 2: enp0s31f6: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500
    qdisc pfifo_fast state UP group default qlen 1000
    link/ether 4c:52:62:0e:54:8b brd ff:ff:ff:ff:ff:ff
8     inet 141.62.66.5/24 brd 141.62.66.255 scope global
        dynamic enp0s31f6
        valid_lft 11902sec preferred_lft 11902sec

```

Senden Sie einen ping-command an einen zweiten Rechner, der am gleichen Switch angeschlossen ist

Hier wird ein anderer Laborrechner, 141.62.66.4, angepingt.

```

$ ping 141.62.66.4
2 PING 141.62.66.4 (141.62.66.4) 56(84) bytes of data.
    64 bytes from 141.62.66.4: icmp_seq=1 ttl=64 time=0.670 ms

```




Figure 6: Der gegenüberliegende Netzwerkschrank wird durch Glasfaser erreicht. Wie im Bild zu sehen, sind zwei Glasfaserkabel an das Panel mit der Aufschrift "Panel B" angeschlossen. Zwei Kabel daher, da eines der beiden Kabel für das eingehende Signal reserviert ist und das andere für das ausgehende Signal. Durch diese beiden Kabel sind die Netzwerkschränke miteinander verbunden. Bei Glasfaserkabel muss beachtet werden, dass die Kabel nicht zu stark gebogen sind, da dies sonst zu Signalverlust führt.


```

4 64 bytes from 141.62.66.4: icmp_seq=2 ttl=64 time=0.509 ms
  64 bytes from 141.62.66.4: icmp_seq=3 ttl=64 time=0.532 ms
6 64 bytes from 141.62.66.4: icmp_seq=4 ttl=64 time=0.526 ms
  64 bytes from 141.62.66.4: icmp_seq=5 ttl=64 time=0.533 ms
8 ^C
  — 141.62.66.4 ping statistics —
10 5 packets transmitted, 5 received, 0% packet loss, time 4085ms
   rtt min/avg/max/mdev = 0.509/0.554/0.670/0.058 ms

```

Senden Sie einen ping-command zu einem Rechner, der am Switch im gegenüberliegenden Netzwerkschrank angeschlossen ist

Hier wird nun ein Rechner mit der IP 141.62.66.13 angepingt, welcher am Switch im gegenüberliegenden Netzwerkschrank angeschlossen ist. Wie zu sehen ist ist die Latenz um ~0.2 ms größer.

```

$ ping 141.62.66.13
2 PING 141.62.66.13 (141.62.66.13) 56(84) bytes of data.
  64 bytes from 141.62.66.13: icmp_seq=1 ttl=128 time=0.786 ms
4 64 bytes from 141.62.66.13: icmp_seq=2 ttl=128 time=0.775 ms
  64 bytes from 141.62.66.13: icmp_seq=3 ttl=128 time=0.853 ms
6 64 bytes from 141.62.66.13: icmp_seq=4 ttl=128 time=0.752 ms
  64 bytes from 141.62.66.13: icmp_seq=5 ttl=128 time=0.793 ms
8 ^C
  — 141.62.66.13 ping statistics —
10 5 packets transmitted, 5 received, 0% packet loss, time 4095ms
   rtt min/avg/max/mdev = 0.752/0.791/0.853/0.033 ms

```

Senden Sie einen ping-command zum Labor-Router

Der Labor-Router hat die IP-Adresse 141.62.66.250. Die Latenz beläuft sich bei diesem mal auf ~1.05 ms.

```

$ ping 141.62.66.250
2 PING 141.62.66.250 (141.62.66.250) 56(84) bytes of data.
  64 bytes from 141.62.66.250: icmp_seq=1 ttl=64 time=1.13 ms
4 64 bytes from 141.62.66.250: icmp_seq=2 ttl=64 time=1.07 ms
  64 bytes from 141.62.66.250: icmp_seq=3 ttl=64 time=1.03 ms
6 64 bytes from 141.62.66.250: icmp_seq=4 ttl=64 time=1.02 ms
  64 bytes from 141.62.66.250: icmp_seq=5 ttl=64 time=1.02 ms
8 64 bytes from 141.62.66.250: icmp_seq=6 ttl=64 time=1.03 ms
  ^C
10 — 141.62.66.250 ping statistics —
   6 packets transmitted, 6 received, 0% packet loss, time 5007ms
12 rtt min/avg/max/mdev = 1.015/1.046/1.127/0.040 ms

```

Starten Sie einen Web-Browser und überprüfen Sie die korrekte Funktion des DNS-Servers durch Aufruf einer beliebigen URL

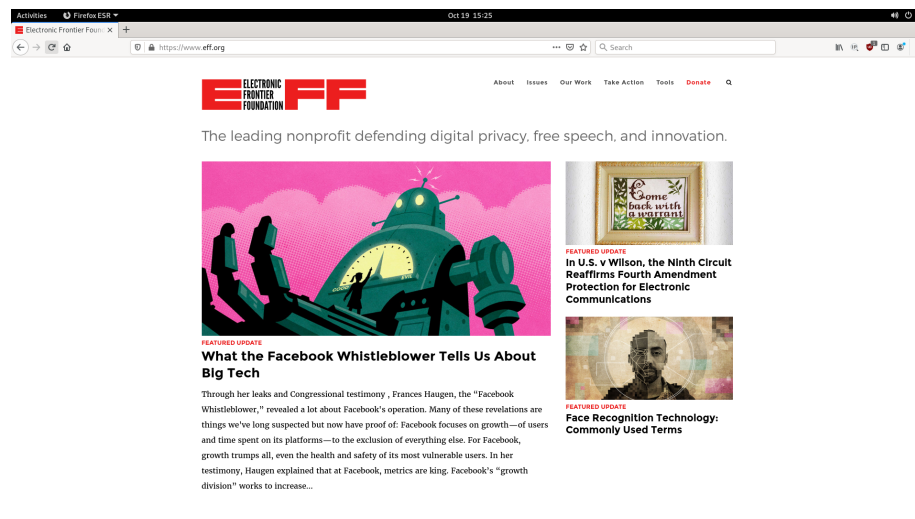


Figure 7: Screenshot

Die Seite ist erreichbar und war davor nicht gecached. Daraus lässt sich schließen, dass die DNS-Abfrage erfolgreich funktioniert hat.

Sehen Sie sich den DNS-Cache an

```
$ sudo journalctl -u systemd-resolved
2 — Journal begins at Tue 2021-10-05 07:59:05 CEST, ends at
   Tue 2021-10-19 15:33:33 CEST. —
   Oct 19 15:31:00 rn05 systemd[1]: Starting Network Name
   Resolution...
4 Oct 19 15:31:00 rn05 systemd-resolved[34579]: Positive Trust
   Anchors:
   Oct 19 15:31:00 rn05 systemd-resolved[34579]: . IN DS 20326 8
   2
   e06d44b80b8f1d39a95c0b0d7c65d08458e880409bbc683457104237c7f8ec8d
6 Oct 19 15:31:00 rn05 systemd-resolved[34579]: Negative trust
   anchors: 10.in-addr.arpa 16.172.in-addr.arpa
   17.172.in-addr.arpa 18.172.in-addr.arpa
   19.172.in-addr.arpa 20.172.in-addr.arpa
   21.172.in-addr.arpa 22.172.in-addr.arpa
   23.172.in-addr.arpa 24.172.in-addr.arpa
   25.172.in-addr.arpa 26.172.in-addr.arpa
   27.172.in-addr.arpa 28.172.in-addr.arpa
   29.172.in-addr.arpa 30.172.in-addr.arpa
   31.172.in-addr.arpa 168.192.in-addr.arpa d.f.ip6.arpa
   corp home internal intranet lan local private test
```

```

Oct 19 15:31:00 rn05 systemd-resolved[34579]: Using system
hostname 'rn05'.
8 Oct 19 15:31:00 rn05 systemd[1]: Started Network Name
Resolution.
Oct 19 15:31:29 rn05 systemd-resolved[34579]: [Scope
protocol=llmnr interface=enp0s31f6 family=AF_INET]
10 Oct 19 15:31:29 rn05 systemd-resolved[34579]: ZONE:
Oct 19 15:31:29 rn05 systemd-resolved[34579]:
5.66.62.141.in-addr.arpa IN PTR rn05
12 Oct 19 15:31:29 rn05 systemd-resolved[34579]: rn05 IN
A 141.62.66.5
Oct 19 15:31:29 rn05 systemd-resolved[34579]: [Scope
protocol=dns]
14 Oct 19 15:31:29 rn05 systemd-resolved[34579]: [Server
141.62.66.250 type=system]
Oct 19 15:31:29 rn05 systemd-resolved[34579]:
Verified feature level: n/a
16 Oct 19 15:31:29 rn05 systemd-resolved[34579]:
Possible feature level: TLS+EDNS0+D0
Oct 19 15:31:29 rn05 systemd-resolved[34579]: DNSSEC
Mode: no
18 Oct 19 15:31:29 rn05 systemd-resolved[34579]: Can do
DNSSEC: yes
Oct 19 15:31:29 rn05 systemd-resolved[34579]: Maximum
UDP packet size received: 512
20 Oct 19 15:31:29 rn05 systemd-resolved[34579]: Failed
UDP attempts: 0
Oct 19 15:31:29 rn05 systemd-resolved[34579]: Failed
TCP attempts: 0
22 Oct 19 15:31:29 rn05 systemd-resolved[34579]: Seen
truncated packet: no
Oct 19 15:31:29 rn05 systemd-resolved[34579]: Seen
OPT RR getting lost: no
24 Oct 19 15:31:29 rn05 systemd-resolved[34579]: Seen
RRSIG RR missing: no
Oct 19 15:32:38 rn05 systemd-resolved[34579]: [Scope
protocol=llmnr interface=enp0s31f6 family=AF_INET]
26 Oct 19 15:32:38 rn05 systemd-resolved[34579]: ZONE:
Oct 19 15:32:38 rn05 systemd-resolved[34579]:
5.66.62.141.in-addr.arpa IN PTR rn05
28 Oct 19 15:32:38 rn05 systemd-resolved[34579]: rn05 IN
A 141.62.66.5
Oct 19 15:32:38 rn05 systemd-resolved[34579]: [Scope
protocol=dns]
30 Oct 19 15:32:38 rn05 systemd-resolved[34579]: [Server
141.62.66.250 type=system]
Oct 19 15:32:38 rn05 systemd-resolved[34579]:
Verified feature level: n/a
32 Oct 19 15:32:38 rn05 systemd-resolved[34579]:
Possible feature level: TLS+EDNS0+D0

```

```

Oct 19 15:32:38 rn05 systemd-resolved[34579]: DNSSEC
Mode: no
34 Oct 19 15:32:38 rn05 systemd-resolved[34579]: Can do
DNSSEC: yes
Oct 19 15:32:38 rn05 systemd-resolved[34579]: Maximum
UDP packet size received: 512
36 Oct 19 15:32:38 rn05 systemd-resolved[34579]: Failed
UDP attempts: 0
Oct 19 15:32:38 rn05 systemd-resolved[34579]: Failed
TCP attempts: 0
38 Oct 19 15:32:38 rn05 systemd-resolved[34579]: Seen
truncated packet: no
Oct 19 15:32:38 rn05 systemd-resolved[34579]: Seen
OPT RR getting lost: no
40 Oct 19 15:32:38 rn05 systemd-resolved[34579]: Seen
RRSIG RR missing: no
Oct 19 15:33:00 rn05 systemd-resolved[34579]: [Scope
protocol=llmnr interface=enp0s31f6 family=AF_INET]
42 Oct 19 15:33:00 rn05 systemd-resolved[34579]: ZONE:
Oct 19 15:33:00 rn05 systemd-resolved[34579]:
5.66.62.141.in-addr.arpa IN PTR rn05
44 Oct 19 15:33:00 rn05 systemd-resolved[34579]: rn05 IN
A 141.62.66.5
Oct 19 15:33:00 rn05 systemd-resolved[34579]: [Scope
protocol=dns]
46 Oct 19 15:33:00 rn05 systemd-resolved[34579]: CACHE:
Oct 19 15:33:00 rn05 systemd-resolved[34579]:
test.com IN A 67.225.146.248
48 Oct 19 15:33:00 rn05 systemd-resolved[34579]:
test.com IN AAAA — NODATA
Oct 19 15:33:00 rn05 systemd-resolved[34579]: [Server
141.62.66.250 type=system]
50 Oct 19 15:33:00 rn05 systemd-resolved[34579]:
Verified feature level: UDP+EDNS0
Oct 19 15:33:00 rn05 systemd-resolved[34579]:
Possible feature level: UDP+EDNS0
52 Oct 19 15:33:00 rn05 systemd-resolved[34579]: DNSSEC
Mode: no
Oct 19 15:33:00 rn05 systemd-resolved[34579]: Can do
DNSSEC: no
54 Oct 19 15:33:00 rn05 systemd-resolved[34579]: Maximum
UDP packet size received: 512
Oct 19 15:33:00 rn05 systemd-resolved[34579]: Failed
UDP attempts: 0
56 Oct 19 15:33:00 rn05 systemd-resolved[34579]: Failed
TCP attempts: 0
Oct 19 15:33:00 rn05 systemd-resolved[34579]: Seen
truncated packet: no
58 Oct 19 15:33:00 rn05 systemd-resolved[34579]: Seen
OPT RR getting lost: no

```

```

Oct 19 15:33:00 rn05 systemd-resolved[34579]: Seen
RRSIG RR missing: no
60 Oct 19 15:33:30 rn05 systemd-resolved[34579]: [Scope
    protocol=llmnr interface=enp0s31f6 family=AF_INET]
Oct 19 15:33:30 rn05 systemd-resolved[34579]: ZONE:
62 Oct 19 15:33:30 rn05 systemd-resolved[34579]:
    5.66.62.141.in-addr.arpa IN PTR rn05
Oct 19 15:33:30 rn05 systemd-resolved[34579]: rn05 IN
    A 141.62.66.5
64 Oct 19 15:33:30 rn05 systemd-resolved[34579]: [Scope
    protocol=dns]
Oct 19 15:33:30 rn05 systemd-resolved[34579]: CACHE:
66 Oct 19 15:33:30 rn05 systemd-resolved[34579]:
    test.com IN AAAA — NODATA
Oct 19 15:33:30 rn05 systemd-resolved[34579]:
    example.com IN AAAA 2606:2800:220:1:248:1893:25c8:1946
68 Oct 19 15:33:30 rn05 systemd-resolved[34579]:
    test.com IN A 67.225.146.248
Oct 19 15:33:30 rn05 systemd-resolved[34579]:
    example.com IN A 93.184.216.34
70 Oct 19 15:33:30 rn05 systemd-resolved[34579]: [Server
    141.62.66.250 type=system]
Oct 19 15:33:30 rn05 systemd-resolved[34579]:
    Verified feature level: UDP+EDNS0
72 Oct 19 15:33:30 rn05 systemd-resolved[34579]:
    Possible feature level: UDP+EDNS0
Oct 19 15:33:30 rn05 systemd-resolved[34579]: DNSSEC
    Mode: no
74 Oct 19 15:33:30 rn05 systemd-resolved[34579]: Can do
    DNSSEC: no
Oct 19 15:33:30 rn05 systemd-resolved[34579]: Maximum
    UDP packet size received: 512
76 Oct 19 15:33:30 rn05 systemd-resolved[34579]: Failed
    UDP attempts: 0
Oct 19 15:33:30 rn05 systemd-resolved[34579]: Failed
    TCP attempts: 0
78 Oct 19 15:33:30 rn05 systemd-resolved[34579]: Seen
    truncated packet: no
Oct 19 15:33:30 rn05 systemd-resolved[34579]: Seen
    OPT RR getting lost: no
80 Oct 19 15:33:30 rn05 systemd-resolved[34579]: Seen
    RRSIG RR missing: no

```

Wie zu erkennen ist, befinden sich mom. 2 Einträge im DNS-Cache: test.com und example.com, für welche jeweils die A und AAAA-Records gecached wurden.

Adress Resolution Protocol ARP

arp ist deprecated, es wird stattdessen ip neigh verwendet.

Dokumentieren Sie den Inhalt der ARP-Tabelle Ihres PC (arp-a, DOS-Box).

```
$ ip neigh show
2 141.62.66.186 dev enp0s31f6 lladdr 10:82:86:01:36:6d STALE
  141.62.66.12 dev enp0s31f6 lladdr 4c:52:62:0e:e0:e9 STALE
4 141.62.66.14 dev enp0s31f6 lladdr 4c:52:62:0e:e0:ae STALE
  141.62.66.250 dev enp0s31f6 lladdr 00:0d:b9:4f:b8:14 REACHABLE
6 141.62.66.4 dev enp0s31f6 lladdr 4c:52:62:0e:53:eb STALE
  141.62.66.13 dev enp0s31f6 lladdr 4c:52:62:0e:54:5d STALE
8 141.62.66.22 dev enp0s31f6 FAILED
  141.62.66.216 dev enp0s31f6 lladdr 44:31:92:50:6c:61 STALE
```

Nun pingen Sie einen beliebigen anderen Arbeitsplatz an und beobachten Sie evtl. Veränderungen der ARP-Tabelle

```
1 $ ping 141.62.66.236
PING 141.62.66.236 (141.62.66.236) 56(84) bytes of data.
3 64 bytes from 141.62.66.236: icmp_seq=1 ttl=64 time=0.530 ms
  64 bytes from 141.62.66.236: icmp_seq=2 ttl=64 time=0.684 ms
5 64 bytes from 141.62.66.236: icmp_seq=3 ttl=64 time=0.424 ms
^C
7 — 141.62.66.236 ping statistics —
  3 packets transmitted, 3 received, 0% packet loss, time 2031ms
9 $ ip neigh show
  141.62.66.186 dev enp0s31f6 lladdr 10:82:86:01:36:6d STALE
11 141.62.66.12 dev enp0s31f6 lladdr 4c:52:62:0e:e0:e9 STALE
  141.62.66.236 dev enp0s31f6 lladdr 26:c5:04:8a:fa:eb STALE
13 141.62.66.14 dev enp0s31f6 lladdr 4c:52:62:0e:e0:ae STALE
  141.62.66.250 dev enp0s31f6 lladdr 00:0d:b9:4f:b8:14 REACHABLE
15 141.62.66.4 dev enp0s31f6 lladdr 4c:52:62:0e:53:eb STALE
  141.62.66.13 dev enp0s31f6 lladdr 4c:52:62:0e:54:5d STALE
17 141.62.66.22 dev enp0s31f6 FAILED
  141.62.66.216 dev enp0s31f6 lladdr 44:31:92:50:6c:61 STALE
```

Nun wurde die Adresse 141.62.66.236 zur ARP-Tabelle hinzugefügt.

Ist die MAC-Adresse Ihres PC lokal oder global vergeben?

```
$ ip a
2 1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state
   UNKNOWN group default qlen 1000
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
4   inet 127.0.0.1/8 scope host lo
      valid_lft forever preferred_lft forever
6 2: enp0s31f6: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500
   qdisc pfifo_fast state UP group default qlen 1000
   link/ether 4c:52:62:0e:54:8b brd ff:ff:ff:ff:ff:ff
8   inet 141.62.66.5/24 brd 141.62.66.255 scope global
      dynamic enp0s31f6
      valid_lft 10201sec preferred_lft 10201sec
```

Es findet sich die MAC-Adresse 4c:52:62:0e:54:8b; ein Lookup der OUI ergibt: 4C:52:62 Fujitsu Technology Solutions GmbH, woraus sich schließen lässt, dass die MAC global vergeben ist.

Was würde geschehen, wenn ein weiterer PC mit gleicher IP (aber selbstverständlich anderer MAC) ans gleiche Subnetz angeschlossen würde?

Ein reines Ethernet-Frame würde den Host noch korrekt erreichen, aber da die IP nun mehreren Hosts zugeordnet wäre würden IP-Pakete nicht mehr den richtigen Host erreichen.

Vergleichen Sie die Vorteile / Nachteile einer statischen und dynamische ARP-Tabelle

Vorteile einer statischen/Nachteile einer dynamischen:

- Schneller und weniger Traffic; ARP-Request muss nicht gemacht werden
- Chain of Trust ist kürzer, da nicht dem Host, welche den ARP-Request beantwortet, vertraut werden muss

Vorteile einer dynamischen/Nachteile einer statischen:

- Wenn Geräte entfernt werden, dann müssen die Einträge manuell gelöscht werden
- Neue Geräte müssen nicht manuell hinzugefügt werden

Warum wird die ARP-Tabelle ganz oder teilweise nach Ablauf einer bestimmten Zeit gelöscht, wie Sie leicht nachvollziehen können?

Durch die Löschung der ARP-Tabelle werden die ARP-Anfragen erneut gemacht; wenn Geräte zum Netzwerk hinzukommen oder entfernt werden, so werden diese Änderungen dadurch repräsentiert.

Ping

Ping-Nutzung

```
$ ping --help
2 Usage
  ping [options] <destination>
4
  Options:
6  <destination>      dns name or ip address
   -a                  use audible ping
8  -A                  use adaptive ping
   -B                  sticky source address
10 -c <count>          stop after <count> replies
```

```

-D          print timestamps
12  -d          use SO_DEBUG socket option
      -f          flood ping
14  -h          print help and exit
      -I <interface> either interface name or address
16  -i <interval> seconds between sending each packet
      -L          suppress loopback of multicast packets
18  -l <preload> send <preload> number of packages while
      waiting replies
      -m <mark>    tag the packets going out
20  -M <pmtud opt> define mtu discovery, can be one of
      <do|dont|want>
      -n          no dns name resolution
22  -O          report outstanding replies
      -p <pattern> contents of padding byte
24  -q          quiet output
      -Q <tclass>  use quality of service <tclass> bits
26  -s <size>    use <size> as number of data bytes to be
      sent
      -S <size>    use <size> as SO_SNDBUF socket option
      value
28  -t <ttl>     define time to live
      -U          print user-to-user latency
30  -v          verbose output
      -V          print version and exit
32  -w <deadline> reply wait <deadline> in seconds
      -W <timeout> time to wait for response
34
IPv4 options:
36  -4          use IPv4
      -b          allow pinging broadcast
38  -R          record route
      -T <timestamp> define timestamp, can be one of
      <tsonly|tsandaddr|tsprespec>
40
IPv6 options:
42  -6          use IPv6
      -F <flowlabel> define flow label, default is random
44  -N <nodeinfo opt> use icmp6 node info query, try <help> as
      argument

```

46 For more details see ping(8).

Erzwungenes IPv4:

```

$ ping -4 google.com
2 PING google.com (142.250.185.78) 56(84) bytes of data:
  64 bytes from fra16s48-in-f14.1e100.net (142.250.185.78):
    icmp_seq=1 ttl=114 time=4.58 ms
4 64 bytes from fra16s48-in-f14.1e100.net (142.250.185.78):

```



```

        icmp_seq=2 ttl=114 time=5.40 ms
^C
6 — google.com ping statistics —
  2 packets transmitted, 2 received, 0% packet loss, time 1002ms
8 rtt min/avg/max/mdev = 4.582/4.989/5.397/0.407 ms

```

Nur zwei Pakete:

```

praktikum@rn05:~$ ping -c 2 google.com
2 PING google.com (142.250.185.78) 56(84) bytes of data.
  64 bytes from fra16s48-in-f14.1e100.net (142.250.185.78):
    icmp_seq=1 ttl=114 time=4.45 ms
4 64 bytes from fra16s48-in-f14.1e100.net (142.250.185.78):
    icmp_seq=2 ttl=114 time=4.46 ms

6 — google.com ping statistics —
  2 packets transmitted, 2 received, 0% packet loss, time 1002ms
8 rtt min/avg/max/mdev = 4.447/4.453/4.460/0.006 ms

```

2 Sekunden Pause zwischen den Paketen:

```

$ ping -i 2 google.com
2 PING google.com (142.250.185.78) 56(84) bytes of data.
  64 bytes from fra16s48-in-f14.1e100.net (142.250.185.78):
    icmp_seq=1 ttl=114 time=4.69 ms
4 64 bytes from fra16s48-in-f14.1e100.net (142.250.185.78):
    icmp_seq=2 ttl=114 time=4.59 ms
^C
6 — google.com ping statistics —
  2 packets transmitted, 2 received, 0% packet loss, time 2003ms
8 rtt min/avg/max/mdev = 4.586/4.639/4.693/0.053 ms

```

HRPing-Nutzung

HRPing ist ein erweitertes Ping-Command mit folgenden Optionen:

```

$ wine64 hrping.exe
2 This is hrPING v5.04 by cFos Software GmbH —
  http://www.cfos.de

4 usage: hrPING [options] host

6 data options:
    -f                Set Don't Fragment bit in IP header
8  -i TTL             Time To Live (default 255 for ping, 30 for
                      traceroute)
    -v TOS            Type Of Service (default 0, deprecated)
10 -l size            Send buffer size (payload size, default 32)
    -l s1[:s2[:i]]    Size sweep: send buffer size from <s1> to
                      <s2> step <i>
12 -L s1[:s2[:i]]    IP datagram size (payload size + 28,
                      default 60) [with sweep]

```

```

-M          Send ICMP timestamp requests
14  -u [port] Send UDP packets (port 7 by default)

16 operational options:
    -t          Ping the specified host until stopped (Ctrl-C
        to stop)
18  -n count    Number of packets to send (default 4)
    -w timeout  Timeout in msec to wait for a reply (default
        2000)
20  -s time     Sending interval between packets in msec
        (default 500)
    -c [num]    Concurrent sending of up to <num> pings at a
        time (default 1)
22  -r [count]  Be a traceroute (do <count> pings each hop,
        default 3)
    -a [hop]    Resolve addresses to names for traceroute
        (start at <hop>)
24  -p          Trace path to destination, then ping all hops
        on path

26 output options:
    -lic        Show public license and warranty
28  -fwhelp     Print firewall help text
    -F file     Log output into <file> as well, even if -q is
        set
30  -T          Print timestamp in front of each line
    -q[r|e|t]   Be quiet (-qr=no replies, -qe=no errors,
        -qt=no timeouts)
32  -y [sec]    Print summary of the last <sec> secs (default
        10)
    -g -G       Show graph (-gg=close graph on exit, -G use
        running grping.exe)
34  -? -h       This help (-??=more help)

36 hrPING is Freeware, please share it! See www.cfos.de for our
    other solutions:
    — Internet Acceleration via Traffic Shaping      : cFosSpeed
38  — Webserver for home users and professionals     : cFos
    Personal Net
    — IPv6 Connectivity for XP, Vista and Windows 7 : cFos
    IPv6 Link

```

HRPing jedoch ist unfreie Software und respektiert deshalb nicht die digitalen Rechte der Versuchsdurchführenden; zudem funktioniert es nicht auf freien Systemen und der Quellcode steht nicht zur Verfügung, was ein Sicherheitsrisiko darstellt. Stattdessen wurde deshalb die freie Implementation fping verwendet:

```

Name       : fping
2 Version  : 5.0
Release    : 3.fc34

```

```

4 Architecture : x86_64
   Size       : 63 k
6 Source      : fping-5.0-3.fc34.src.rpm
   Repository : @System
8 From repo   : fedora
   Summary    : Scriptable, parallelized ping-like utility
10 URL        : http://www.fping.org/
   License    : BSD with advertising
12 Description : fping is a ping-like program which can
                 determine the
                   : accessibility of multiple hosts using ICMP
                   : echo requests. fping
14             : is designed for parallelized monitoring of
                   : large numbers of
                   : systems, and is developed with ease of use in
                   : scripting in mind.

```

Diese hat ähnliche Optionen:

```

1 $ fping --help
   Usage: fping [options] [targets...]
3
   Probing options:
5   -4, --ipv4           only ping IPv4 addresses
   -6, --ipv6           only ping IPv6 addresses
7   -b, --size=BYTES    amount of ping data to send, in bytes
                       (default: 56)
   -B, --backoff=N      set exponential backoff factor to N
                       (default: 1.5)
9   -c, --count=N       count mode: send N pings to each target
   -f, --file=FILE      read list of targets from a file ( -
                       means stdin)
11  -g, --generate       generate target list (only if no -f
                       specified)
                       (give start and end IP in the target
13                       list, or a CIDR address)
                       (ex. fping -g 192.168.1.0 192.168.1.255
                       or fping -g 192.168.1.0/24)
   -H, --ttl=N          set the IP TTL value (Time To Live hops)
15  -I, --iface=IFACE    bind to a particular interface
   -l, --loop           loop mode: send pings forever
17  -m, --all            use all IPs of provided hostnames (e.g.
                       IPv4 and IPv6), use with -A
   -M, --dontfrag       set the Don't Fragment flag
19  -O, --tos=N          set the type of service (tos) flag on
                       the ICMP packets
   -p, --period=MSEC    interval between ping packets to one
                       target (in ms)
21                       (in loop and count modes, default: 1000
                       ms)

```

```

-r, --retry=N      number of retries (default: 3)
23 -R, --random      random packet data (to foil link data
                    compression)
-S, --src=IP       set source address
25 -t, --timeout=MSEC individual target initial timeout
                    (default: 500 ms,
                    except with -l/-c/-C, where it's the -p
                    period up to 2000 ms)
27
    Output options:
29 -a, --alive       show targets that are alive
-A, --addr         show targets by address
31 -C, --vcount=N   same as -c, report results in verbose
                    format
-D, --timestamp    print timestamp before each output line
33 -e, --elapsed    show elapsed time on return packets
-i, --interval=MSEC interval between sending ping packets
                    (default: 10 ms)
35 -n, --name       show targets by name (-d is equivalent)
-N, --netdata      output compatible for netdata (-l -Q
                    are required)
37 -o, --outage     show the accumulated outage time (lost
                    packets * packet interval)
-q, --quiet        quiet (don't show per-target/per-ping
                    results)
39 -Q, --quiet=SECS same as -q, but show summary every n
                    seconds
-s, --stats        print final stats
41 -u, --unreach    show targets that are unreachable
-v, --version      show version
43 -x, --reachable=N shows if >=N hosts are reachable or not

```

Die Verwendung ist ähnlich wie ping.

Weisen Sie mit Hilfe von HRPING nach, dass ein Ping, der zuerst eine ARP-Auflösung erforderlich macht, zu deutlich erhöhten Antwortzeiten führt.

```

$ fping -e 10.60.43.50
2 10.60.43.50 is alive (70.9 ms)
$ sudo ip -s -s neigh flush all
4 10.60.63.252 dev wlp0s20f3 lladdr 3c:fd:fe:b6:ed:2d ref 1
  used 10/10/10 probes 4 REACHABLE
  10.60.43.50 dev wlp0s20f3 lladdr 7a:11:bd:7c:f9:ff ref 1 used
  2/19/2 probes 4 DELAY
6
*** Round 1, deleting 2 entries ***
8 *** Flush is complete after 1 round ***
$ fping -e 10.60.43.50
10 10.60.43.50 is alive (212 ms)

```

Zu erkennen ist, dass nach der Löschen der ARP-Tabelle eine deutlich längere Antwortzeit zu messen ist.

Traceroute & MTR

Versuchen Sie, den zentralen Peering-Point (DE-CIX) in Deutschland geographisch anhand des Namens zu lokalisieren.

```
$ traceroute de-cix.net
2 traceroute to de-cix.net (46.31.121.136), 30 hops max, 60
  byte packets
  1 opnsense-router.rnlabor.hdm-stuttgart.de (141.62.66.250)
    0.509 ms  1.566 ms  0.991 ms
4 2 ciscovlgw318.hdm-stuttgart.de (141.62.31.246)  2.047 ms
    1.295 ms  1.019 ms
  3 firewall-h.hdm-stuttgart.de (141.62.1.1)  1.118 ms  1.450
    ms  1.120 ms
6 4 * * *
  5 stu-al30-1-te0-0-0-17.belwue.net (129.143.56.53)  3.625
    ms  3.191 ms  3.331 ms
8 6 stu-nwz-a99-hu0-3-0-5.belwue.net (129.143.56.106)  3.030
    ms  1.325 ms  1.440 ms
  7 fra-decix-1-hu0-0-0-4.belwue.net (129.143.60.113)  5.149
    ms fra-decix-1-hu0-0-0-3.belwue.net (129.143.57.127)
    5.283 ms  5.465 ms
10 8 sgw2-te-0-0-2-3-ixp.fra.de-cix.net (80.81.194.116)  7.276
    ms  7.181 ms  7.103 ms
  9 * * *
12 10 * * *
  11 * * *
14 12 * * *
  13 * * *
16 14 *^C
```

1. opnsense-router.rnlabor.hdm-stuttgart.de: Gateway des RN-Labors
2. ciscovlgw318.hdm-stuttgart.de: Gateway zwischen RN-Labor-Router und Firewall
3. firewall-h.hdm-stuttgart.de: Firewall der HdM
4. stu-al30-1-te0-0-0-17.belwue.net und stu-nwz-a99-hu0-3-0-5.belwue.net: Router Belwue in Stuttgart
5. fra-decix-1-hu0-0-0-4.belwue.net: Router Belwue in Frankfurt
6. sgw2-te-0-0-2-3-ixp.fra.de-cix.net: Router DE-CIX in Frankfurt

Zeichnen Sie den Weg eines Pakets zu www.aol.com auf.

```
$ traceroute www.aol.com
2 traceroute to www.aol.com (212.82.100.163), 30 hops max, 60
  byte packets
```

```

1  opnsense.rnlabor.hdm-stuttgart.de (141.62.66.250)  1.284
   ms 0.653 ms 0.956 ms
4  2  ciscovlgw318.hdm-stuttgart.de (141.62.31.246)  1.168 ms
   1.601 ms 2.339 ms
   3  firewall-h.hdm-stuttgart.de (141.62.1.1)  1.800 ms 1.896
   ms 2.378 ms
6  4  * * *
   5  stu-al30-1-te0-0-0-17.belwue.net (129.143.56.53)  3.143
   ms 3.819 ms 3.212 ms
8  6  stu-nwz-a99-hu0-3-0-5.belwue.net (129.143.56.106)  3.510
   ms 2.147 ms 3.579 ms
   7  fra-decix-1-hu0-0-0-3.belwue.net (129.143.57.127)  5.073
   ms 5.193 ms 4.812 ms
10 8  ge-1-3-0.pat1.dee.yahoo.com (80.81.192.115)  5.630 ms
   5.656 ms 5.699 ms
   9  ae-3.pat1.frz.yahoo.com (209.191.112.17)  13.928 ms
   14.322 ms 13.942 ms
12 10 ae-2.pat1.iry.yahoo.com (209.191.112.54)  30.229 ms
   30.613 ms 30.790 ms
   11 et-1-1-2.msrl.ir2.yahoo.com (66.196.65.19)  30.763 ms
   29.649 ms 29.854 ms
14 12 lo0.fab2-1-gdc.ir2.yahoo.com (77.238.190.3)  29.678 ms
   lo0.fab3-1-gdc.ir2.yahoo.com (77.238.190.4)  29.709 ms
   lo0.fab2-1-gdc.ir2.yahoo.com (77.238.190.3)  29.842 ms
   13 usw2-1-lba.ir2.yahoo.com (77.238.190.103)  29.724 ms
   29.602 ms usw1-1-lba.ir2.yahoo.com (77.238.190.102)
   29.750 ms
16 14 media-router-aol71.prod.media.vip.ir2.yahoo.com
   (212.82.100.163) 29.546 ms 30.166 ms 29.797 ms

```

Beobachten Sie Zeitüberschreitungen? Wie können Sie `tracert` so manipulieren, dass möglichst selten Zeitüberschreitungen auftauchen?

Eine Zeitüberschreitung kann zwischen `firewall-h.hdm-stuttgart.de` und `stu-al30-1-te0-0-0-17.belwue.net` erkannt werden; hier wurde versucht das Timeout auf 5 Sekunden mittels `-w` zu setzen und mit `-I` über die Raw Sockets API direkt die Pakete am Kernel-Stack vorbei zu schicken, was jedoch in beiden Fällen die durch `* * *` gekennzeichneten Timeouts nicht umgehen kann.

```
$ traceroute --help
```

```
2 Usage:
```

```

traceroute [ -4dFITnreAUDV ] [ -f first_ttl ] [ -g
gate,... ] [ -i device ] [ -m max_ttl ] [ -N squeries ]
[ -p port ] [ -t tos ] [ -l flow_label ] [ -w
MAX,HERE,NEAR ] [ -q nqueries ] [ -s src_addr ] [ -z
sendwait ] [ --fwmark=num ] host [ packetlen ]

```

```
4 Options:
```

```

-4                               Use IPv4
6  -6                             Use IPv6
   -d  --debug                     Enable socket level debugging

```

```

8  -F  --dont-fragment          Do not fragment packets
   -f first_ttl  --first=first_ttl
10                                Start from the first_ttl hop
                                (instead from 1)
   -g gate,...  --gateway=gate,...
12                                Route packets through the
                                specified gateway
                                (maximum 8 for IPv4 and 127 for
                                IPv6)
14  -I  --icmp                  Use ICMP ECHO for tracerouting
   -T  --tcp                    Use TCP SYN for tracerouting
                                (default port is 80)
16  -i device  --interface=device
                                Specify a network interface to
                                operate with
18  -m max_ttl  --max-hops=max_ttl
                                Set the max number of hops (max
                                TTL to be
20                                reached). Default is 30
   -N squeries  --sim-queries=squeries
22                                Set the number of probes to be
                                tried
                                simultaneously (default is 16)
24  -n                                Do not resolve IP addresses to
                                their domain names
   -p port  --port=port
                                Set the destination port to
                                use. It is either
26                                initial udp port value for
                                "default" method
                                (incremented by each probe,
                                default is 33434), or
28                                initial seq for "icmp"
                                (incremented as well,
                                default from 1), or some
                                constant destination
30                                port for other methods (with
                                default of 80 for
                                "tcp", 53 for "udp", etc.)
32  -t tos  --tos=tos
                                Set the TOS (IPv4 type of
                                service) or TC (IPv6
                                traffic class) value for
                                outgoing packets
34  -l flow_label  --flowlabel=flow_label
                                Use specified flow_label for
                                IPv6 packets
36  -w MAX,HERE,NEAR  --wait=MAX,HERE,NEAR
                                Wait for a probe no more than
                                HERE (default 3)
38                                times longer than a response
                                from the same hop,

```

```

                                or no more than NEAR (default
                                10) times than some
40                                next hop, or MAX (default 5.0)
                                seconds (float
                                point values allowed too)
42  -q nqueries  --queries=nqueries
                                Set the number of probes per
                                each hop. Default is
44                                3
                                Bypass the normal routing and
                                send directly to a
46                                host on an attached network
-s src_addr  --source=src_addr
48                                Use source src_addr for
                                outgoing packets
-z sendwait  --sendwait=sendwait
50                                Minimal time interval between
                                probes (default 0).
                                If the value is more than 10,
                                then it specifies a
52                                number in milliseconds, else it
                                is a number of
                                seconds (float point values
                                allowed too)
54  -e --extensions
                                Show ICMP extensions (if
                                present), including MPLS
-A --as-path-lookups
56                                Perform AS path lookups in
                                routing registries and
                                print results directly after
                                the corresponding
                                addresses
58  -M name  --module=name
                                Use specified module (either
                                builtin or external)
                                for traceroute operations. Most
                                methods have
60                                their shortcuts ('-I' means '-M
                                icmp' etc.)
-O OPTS,...  --options=OPTS,...
62                                Use module-specific option OPTS
                                for the
                                traceroute module. Several OPTS
                                allowed,
64                                separated by comma. If OPTS is
                                "help", print info
                                about available options
66  --sport=num
                                Use source port num for
                                outgoing packets. Implies
                                '-N 1'
68  --fwmark=num
                                Set firewall mark for outgoing
                                packets

```



```

-U  --udp                                Use UDP to particular port for
    tracerouting
70                                     (instead of increasing the port
                                     per each probe),
                                     default port is 53
72  -UL                                  Use UDPLITE for tracerouting
    (default dest port
                                     is 53)
74  -D  --dccp                            Use DCCP Request for
    tracerouting (default port
                                     is 33434)
76  -P  prot  --protocol=prot            Use raw packet of protocol prot
    for tracerouting
    --mtu                                Discover MTU along the path
    being traced. Implies
78                                     '-F -N 1'
    --back                               Guess the number of hops in the
    backward path and
80                                     print if it differs
    -V  --version                        Print version info and exit
82  --help                              Read this help and exit

84 Arguments:
+    host                                The host to traceroute to
86    packetlen                          The full packet length (default is the
    length of an IP
    header plus 40). Can be ignored or
    increased to a minimal
88    allowed value
$ traceroute www.aol.com
90 traceroute to www.aol.com (212.82.100.163), 30 hops max, 60
    byte packets
    1  opnsense.rnlabor.hdm-stuttgart.de (141.62.66.250)  1.284
        ms  0.653 ms  0.956 ms
92  2  ciscovlgw318.hdm-stuttgart.de (141.62.31.246)  1.168 ms
        1.601 ms  2.339 ms
    3  firewall-h.hdm-stuttgart.de (141.62.1.1)  1.800 ms  1.896
        ms  2.378 ms
94  4  * * *
    5  stu-al30-1-te0-0-0-17.belwue.net (129.143.56.53)  3.143
        ms  3.819 ms  3.212 ms
96  6  stu-nwz-a99-hu0-3-0-5.belwue.net (129.143.56.106)  3.510
        ms  2.147 ms  3.579 ms
    7  fra-decix-1-hu0-0-0-3.belwue.net (129.143.57.127)  5.073
        ms  5.193 ms  4.812 ms
98  8  ge-1-3-0.pat1.dee.yahoo.com (80.81.192.115)  5.630 ms
        5.656 ms  5.699 ms
    9  ae-3.pat1.frz.yahoo.com (209.191.112.17)  13.928 ms
        14.322 ms  13.942 ms
100 10 ae-2.pat1.iry.yahoo.com (209.191.112.54)  30.229 ms

```

```

30.613 ms 30.790 ms
11 et-1-1-2.msrl.ir2.yahoo.com (66.196.65.19) 30.763 ms
29.649 ms 29.854 ms
102 12 lo0.fab2-1-gdc.ir2.yahoo.com (77.238.190.3) 29.678 ms
lo0.fab3-1-gdc.ir2.yahoo.com (77.238.190.4) 29.709 ms
lo0.fab2-1-gdc.ir2.yahoo.com (77.238.190.3) 29.842 ms
13 usw2-1-lba.ir2.yahoo.com (77.238.190.103) 29.724 ms
29.602 ms usw1-1-lba.ir2.yahoo.com (77.238.190.102)
29.750 ms
104 14 media-router-aol71.prod.media.vip.ir2.yahoo.com
(212.82.100.163) 29.546 ms 30.166 ms 29.797 ms
[pojntfx@felixs-xps13 hrping-v504]$ ssh
pojntfx@159.223.25.154 "nc -lp 6969"
106 $ traceroute -w 5 www.aol.com
traceroute to www.aol.com (212.82.100.163), 30 hops max, 60
byte packets
108 1 opnsense.rnlabor.hdm-stuttgart.de (141.62.66.250) 0.707
ms 3.001 ms 1.312 ms
2 ciscovlgw318.hdm-stuttgart.de (141.62.31.246) 1.782 ms
2.642 ms 2.615 ms
110 3 firewall-h.hdm-stuttgart.de (141.62.1.1) 3.417 ms 0.907
ms 2.692 ms
4 * * *
112 5 stu-a130-1-te0-0-0-17.belwue.net (129.143.56.53) 2.044
ms 2.630 ms 2.032 ms
6 stu-nwz-a99-hu0-3-0-5.belwue.net (129.143.56.106) 3.323
ms 1.287 ms 1.541 ms
114 7 fra-decix-1-hu0-0-0-4.belwue.net (129.143.60.113) 7.004
ms 7.114 ms 7.266 ms
8 ge-1-3-0.pat1.dee.yahoo.com (80.81.192.115) 6.009 ms
4.880 ms 4.545 ms
116 9 ae-3.pat1.frz.yahoo.com (209.191.112.17) 14.326 ms
13.727 ms 13.700 ms
10 ae-2.pat1.iry.yahoo.com (209.191.112.54) 31.291 ms
31.060 ms 31.097 ms
118 11 ge-0-3-9-d104.pat1.the.yahoo.com (66.196.65.21) 29.823
ms 29.921 ms et-1-1-2.msrl.ir2.yahoo.com (66.196.65.19)
29.735 ms
12 lo0.fab4-1-gdc.ir2.yahoo.com (77.238.190.5) 29.809 ms
lo0.fab1-1-gdc.ir2.yahoo.com (77.238.190.2) 29.664 ms
29.659 ms
120 13 usw1-1-lba.ir2.yahoo.com (77.238.190.102) 29.517 ms
29.572 ms 29.759 ms
14 media-router-aol71.prod.media.vip.ir2.yahoo.com
(212.82.100.163) 29.563 ms 29.706 ms 29.883 ms
122 $ sudo traceroute -I www.aol.com
traceroute to www.aol.com (212.82.100.163), 30 hops max, 60
byte packets
124 1 opnsense-router.rnlabor.hdm-stuttgart.de (141.62.66.250)
0.461 ms 0.551 ms 0.664 ms

```

```

2   ciscovlgw318.hdm-stuttgart.de (141.62.31.246)  2.064 ms
    2.290 ms  2.657 ms
126 3   firewall-h.hdm-stuttgart.de (141.62.1.1)    1.315 ms  1.628
    ms  1.878 ms
4   * * *
128 5   stu-al30-1-te0-0-0-17.belwue.net (129.143.56.53)  2.891
    ms  3.008 ms  3.068 ms
    6   stu-nwz-a99-hu0-3-0-5.belwue.net (129.143.56.106)  3.175
    ms  1.587 ms  1.432 ms
130 7   fra-decix-1-hu0-0-0-3.belwue.net (129.143.57.127)  5.115
    ms  5.213 ms  5.328 ms
    8   ge-1-3-0.pat1.dee.yahoo.com (80.81.192.115)  4.916 ms
    4.915 ms  5.005 ms
132 9   ae-3.pat1.frz.yahoo.com (209.191.112.17)  13.831 ms
    13.886 ms  14.163 ms
    10  ae-2.pat1.iry.yahoo.com (209.191.112.54)  30.506 ms
    30.505 ms  30.108 ms
134 11  ge-0-3-9-d104.pat1.the.yahoo.com (66.196.65.21)  29.434
    ms  29.657 ms  29.699 ms
    12  lo0.fab3-1-gdc.ir2.yahoo.com (77.238.190.4)  29.757 ms
    29.662 ms  29.707 ms
136 13  usw2-1-lba.ir2.yahoo.com (77.238.190.103)  29.685 ms
    29.690 ms  29.696 ms
    14  media-router-aol71.prod.media.vip.ir2.yahoo.com
    (212.82.100.163)  29.631 ms  29.915 ms  30.152 ms

```

Besuchen Sie das DENIC (www.denic.de) und erfragen Sie den Besitzer von Domain-Namen, die Sie interessieren.

Hier z.B. die HdM Stuttgart:

```

1 $ whois www.hdm-stuttgart.de
   [Querying whois.denic.de]
3 [whois.denic.de]
   % Restricted rights.
5 %
   % Terms and Conditions of Use
7 %
   % The above data may only be used within the scope of
   technical or
9 % administrative necessities of Internet operation or to
   remedy legal
   % problems.
11 % The use for other purposes, in particular for advertising,
   is not permitted.
   %
13 % The DENIC whois service on port 43 doesn't disclose any
   information concerning
   % the domain holder, general request and abuse contact.
15 % This information can be obtained through use of our
   web-based whois service

```

% available at the DENIC website:
 17 % <http://www.denic.de/en/domains/whois-service/web-whois.html>
 %
 19 %

21 Domain: hdm-stuttgart.de
 Nserver: dns1.belwue.de
 23 Nserver: dns3.belwue.de
 Nserver: iz-net-2.hdm-stuttgart.de 141.62.1.2
 25 Nserver: iz-net-3.hdm-stuttgart.de 141.62.1.3
 Nserver: iz-net-4.hdm-stuttgart.de 141.62.1.4
 27 Status: connect
 Changed: 2015-04-22T16:37:06+02:00

Und die Electronic Frontier Foundation:

\$ whois eff.org
 2 [Querying whois.pir.org]
 [whois.pir.org]
 4 Domain Name: EFF.ORG
 Registry Domain ID: D2234962-LROR
 6 Registrar WHOIS Server: whois.gandi.net
 Registrar URL: <http://www.gandi.net>
 8 Updated Date: 2018-03-08T02:19:58Z
 Creation Date: 1990-10-10T04:00:00Z
 10 Registry Expiry Date: 2022-10-09T04:00:00Z
 Registrar Registration Expiration Date:
 12 Registrar: Gandi SAS
 Registrar IANA ID: 81
 14 Registrar Abuse Contact Email: abuse@support.gandi.net
 Registrar Abuse Contact Phone: +33.170377661
 16 Reseller:
 Domain Status: clientTransferProhibited
<https://icann.org/epp#clientTransferProhibited>
 18 Registrant Organization: Electronic Frontier Foundation
 Registrant State/Province: CA
 20 Registrant Country: US
 Name Server: NS1.EFF.ORG
 22 Name Server: NS2.EFF.ORG
 Name Server: NS4.EFF.ORG
 24 DNSSEC: unsigned
 URL of the ICANN Whois Inaccuracy Complaint Form
<https://www.icann.org/wicf/>
 26 >>> Last update of WHOIS database: 2021-10-20T20:35:43Z <<<
 28 For more information on Whois status codes, please visit
<https://icann.org/epp>
 30 Access to Public Interest Registry WHOIS information is
 provided to assist persons in determining the contents of

a domain name registration record in the Public Interest Registry registry database. The data in this record is provided by Public Interest Registry for informational purposes only, and Public Interest Registry does not guarantee its accuracy. This service is intended only for query-based access. You agree that you will use this data only for lawful purposes and that, under no circumstances will you use this data to (a) allow, enable, or otherwise support the transmission by e-mail, telephone, or facsimile of mass unsolicited, commercial advertising or solicitations to entities other than the data recipient's own existing customers; or (b) enable high volume, automated, electronic processes that send queries or data to the systems of Registry Operator, a Registrar, or Afilias except as reasonably necessary to register domain names or modify existing registrations. All rights reserved. Public Interest Registry reserves the right to modify these terms at any time. By submitting this query, you agree to abide by this policy.

- 32 The Registrar of Record identified in this output may have an RDDS service that can be queried for additional information on how to contact the Registrant, Admin, or Tech contact of the queried domain name.

Sehen Sie sich die Möglichkeiten von Pathping an.

Als freie Alternative zu Pathping wurde mtr verwendet:

```

Name          : mtr
2 Epoch       : 2
Version       : 0.94
4 Release     : 3.fc34
Architecture  : x86_64
6 Size        : 191 k
Source        : mtr-0.94-3.fc34.src.rpm
8 Repository  : @System
From repo     : updates
10 Summary    : Network diagnostic tool combining 'traceroute'
               and 'ping'
URL           : https://www.bitwizzard.nl/mtr/
12 License    : GPLv2
Description   : MIR combines the functionality of the
               'traceroute' and 'ping'
14            : programs in a single network diagnostic tool.
               :
16            : When MIR is started, it investigates the
               network connection
               : between the host MIR runs on and the
               user-specified destination

```

```

18      : host. Afterwards it determines the address of
      : each network hop
      : between the machines and sends a sequence of
      : ICMP echo requests
20      : to each one to determine the quality of the
      : link to each machine.
      : While doing this, it prints running statistics
      : about each
22      : machine.
      :
24      : MTR provides two user interfaces: an ncurses
      : interface, useful
      : for the command line, e.g. for SSH sessions;
      : and a GTK interface
26      : for X (provided in the mtr-gtk package).

```

mtr kombiniert die Funktionalität von traceroute und ping, was folgende Optionen ermöglicht:

Usage:

```

2  mtr [options] hostname

4  -F, --filename FILE      read hostname(s) from a file
      -4                    use IPv4 only
6  -6                      use IPv6 only
      -u, --udp             use UDP instead of ICMP echo
8  -T, --tcp               use TCP instead of ICMP echo
      -I, --interface NAME  use named network interface
10 -a, --address ADDRESS    bind the outgoing socket to
      ADDRESS
      -f, --first-ttl NUMBER set what TTL to start
12 -m, --max-ttl NUMBER     maximum number of hops
      -U, --max-unknown NUMBER maximum unknown host
14 -P, --port PORT          target port number for TCP, SCTP,
      or UDP
      -L, --localport LOCALPORT source port number for UDP
16 -s, --psize PACKETSIZE   set the packet size used for
      probing
      -B, --bitpattern NUMBER set bit pattern to use in payload
18 -i, --interval SECONDS   ICMP echo request interval
      -G, --gracetime SECONDS number of seconds to wait for
      responses
20 -Q, --tos NUMBER          type of service field in IP header
      -e, --mpls            display information from ICMP
      extensions
22 -Z, --timeout SECONDS    seconds to keep probe sockets open
      -M, --mark MARK       mark each sent packet
24 -r, --report             output using report mode
      -w, --report-wide     output wide report
26 -c, --report-cycles COUNT set the number of pings sent

```

	-j, --json	output json
28	-x, --xml	output xml
	-C, --csv	output comma separated values
30	-l, --raw	output raw format
	-p, --split	split output
32	-t, --curses	use curses terminal interface
	--displaymode MODE	select initial display mode
34	-n, --no-dns	do not resolve host names
	-b, --show-ips	show IP numbers and host names
36	-o, --order FIELDS	select output fields
	-y, --ipinfo NUMBER	select IP information in output
38	-z, --aslookup	display AS number
	-h, --help	display this help and exit
40	-v, --version	output version information and
	exit	

42 See the 'man 8 mtr' for details.

Interessant ist z.B. die -n-Flag:

```
$ mtr -n --json www.aol.com
2 {
4     "report": {
6         "mtr": {
8             "src": "felixs-xps13",
10            "dst": "www.aol.com",
12            "tos": 0,
14            "tests": 10,
16            "psize": "64",
18            "bitpattern": "0x00"
20        },
22        "hubs": [
24            {
26                "count": 1,
28                "host": "10.60.63.252",
30                "Loss%": 0.0,
                "Snt": 10,
                "Last": 88.565,
                "Avg": 10.379,
                "Best": 1.066,
                "Wrst": 88.565,
                "StDev": 27.477
            },
            {
                "count": 2,
                "host": "141.62.31.94",
                "Loss%": 0.0,
                "Snt": 10,
                "Last": 11.83,
                "Avg": 2.541,
```

```

32         "Best ": 1.24 ,
33         "Wrst ": 11.83 ,
34         "StDev ": 3.272
35     },
36     {
37         "count ": 3,
38         "host ": "???",
39         "Loss%": 100.0 ,
40         "Snt ": 10,
41         "Last ": 0.0 ,
42         "Avg ": 0.0 ,
43         "Best ": 0.0 ,
44         "Wrst ": 0.0 ,
45         "StDev ": 0.0
46     },
47     {
48         "count ": 4,
49         "host ": "129.143.56.53",
50         "Loss%": 0.0 ,
51         "Snt ": 10,
52         "Last ": 16.222 ,
53         "Avg ": 3.928 ,
54         "Best ": 1.613 ,
55         "Wrst ": 16.222 ,
56         "StDev ": 4.422
57     },
58     {
59         "count ": 5,
60         "host ": "129.143.56.106",
61         "Loss%": 0.0 ,
62         "Snt ": 10,
63         "Last ": 231.77 ,
64         "Avg ": 25.22 ,
65         "Best ": 1.846 ,
66         "Wrst ": 231.77 ,
67         "StDev ": 72.574
68     },
69     {
70         "count ": 6,
71         "host ": "129.143.60.113",
72         "Loss%": 0.0 ,
73         "Snt ": 10,
74         "Last ": 77.414 ,
75         "Avg ": 13.153 ,
76         "Best ": 5.437 ,
77         "Wrst ": 77.414 ,
78         "StDev ": 22.584
79     },
80     {
81         "count ": 7,

```



```

82         "host ": "80.81.192.115",
        "Loss%": 0.0,
84         "Snt ": 10,
        "Last ": 86.385,
        "Avg ": 13.403,
86         "Best ": 5.122,
        "Wrst ": 86.385,
88         "StDev ": 25.643
    },
90     {
        "count ": 8,
92         "host ": "209.191.112.17",
        "Loss%": 0.0,
94         "Snt ": 10,
        "Last ": 138.72,
96         "Avg ": 29.309,
        "Best ": 13.844,
98         "Wrst ": 138.72,
        "StDev ": 39.424
100    },
102    {
        "count ": 9,
        "host ": "209.191.112.54",
104        "Loss%": 0.0,
        "Snt ": 10,
106        "Last ": 116.04,
        "Avg ": 41.328,
108        "Best ": 29.978,
        "Wrst ": 116.04,
110        "StDev ": 26.988
    },
112    {
        "count ": 10,
114        "host ": "66.196.65.21",
        "Loss%": 0.0,
116        "Snt ": 10,
        "Last ": 39.317,
118        "Avg ": 31.703,
        "Best ": 30.246,
120        "Wrst ": 39.317,
        "StDev ": 2.747
122    },
124    {
        "count ": 11,
        "host ": "77.238.190.5",
126        "Loss%": 0.0,
        "Snt ": 10,
128        "Last ": 32.85,
        "Avg ": 31.768,
130        "Best ": 30.18,

```

```

132         "Wrst ": 38.489,
        "StDev ": 2.535
134     },
    {
136         "count ": 12,
        "host ": "77.238.190.103",
138         "Loss%": 0.0,
        "Snt ": 10,
140         "Last ": 30.614,
        "Avg ": 33.189,
        "Best ": 30.017,
142         "Wrst ": 56.002,
        "StDev ": 8.102
144     },
    {
146         "count ": 13,
        "host ": "212.82.100.163",
148         "Loss%": 0.0,
        "Snt ": 10,
150         "Last ": 32.157,
        "Avg ": 30.531,
152         "Best ": 29.846,
        "Wrst ": 32.157,
154         "StDev ": 0.818
    }
156 ]
}
158 }
$ mtr --json www.aol.com
160 {
    "report": {
162         "mtr": {
            "src ": "felixs-xps13",
164             "dst ": "www.aol.com",
            "tos ": 0,
166             "tests ": 10,
            "psize ": "64",
168             "bitpattern ": "0x00"
        },
170         "hubs": [
            {
172                 "count ": 1,
                "host ": "_gateway",
174                 "Loss%": 0.0,
                "Snt ": 10,
176                 "Last ": 35.643,
                "Avg ": 5.191,
178                 "Best ": 1.074,
                "Wrst ": 35.643,
180                 "StDev ": 10.757
            }
        ]
    }
}

```

```

    },
182  {
    "count ": 2,
184  "host ": "141.62.31.94",
    "Loss%": 0.0,
186  "Snt ": 10,
    "Last ": 49.069,
188  "Avg ": 14.104,
    "Best ": 1.404,
190  "Wrst ": 77.221,
    "StDev": 26.687
192  },
    {
194  "count ": 3,
    "host ": "???",
196  "Loss%": 100.0,
    "Snt ": 10,
198  "Last ": 0.0,
    "Avg ": 0.0,
200  "Best ": 0.0,
    "Wrst ": 0.0,
202  "StDev": 0.0
    },
204  {
    "count ": 4,
206  "host ": "stu-al30-1-te0-0-0-17.belwue.net",
    "Loss%": 0.0,
208  "Snt ": 10,
    "Last ": 14.869,
210  "Avg ": 11.953,
    "Best ": 1.886,
212  "Wrst ": 50.552,
    "StDev": 17.083
214  },
    {
216  "count ": 5,
    "host ": "stu-nwz-a99-hu0-3-0-5.belwue.net",
218  "Loss%": 0.0,
    "Snt ": 10,
220  "Last ": 2.332,
    "Avg ": 2.954,
222  "Best ": 1.847,
    "Wrst ": 7.302,
224  "StDev": 1.961
    },
226  {
    "count ": 6,
228  "host ": "fra-decix-1-hu0-0-0-4.belwue.net",
    "Loss%": 0.0,
230  "Snt ": 10,

```

```

232         "Last ": 58.059,
        "Avg ": 22.657,
234         "Best ": 5.208,
        "Wrst ": 74.371,
        "StDev ": 27.785
236     },
    {
238         "count ": 7,
        "host ": "ge-1-3-0.pat1.dee.yahoo.com",
240         "Loss%": 0.0,
        "Snt ": 10,
242         "Last ": 5.488,
        "Avg ": 6.379,
244         "Best ": 4.908,
        "Wrst ": 13.858,
246         "StDev ": 2.716
    },
248     {
        "count ": 8,
250         "host ": "ae-3.pat1.frz.yahoo.com",
        "Loss%": 0.0,
252         "Snt ": 10,
        "Last ": 125.22,
254         "Avg ": 33.495,
        "Best ": 14.004,
256         "Wrst ": 125.22,
        "StDev ": 40.562
258     },
    {
260         "count ": 9,
        "host ": "ae-2.pat1.iry.yahoo.com",
262         "Loss%": 0.0,
        "Snt ": 10,
264         "Last ": 86.809,
        "Avg ": 36.314,
266         "Best ": 29.889,
        "Wrst ": 86.809,
268         "StDev ": 17.76
    },
270     {
        "count ": 10,
272         "host ": "ge-0-3-9-d104.pat1.the.yahoo.com",
        "Loss%": 0.0,
274         "Snt ": 10,
        "Last ": 31.651,
276         "Avg ": 41.326,
        "Best ": 30.095,
278         "Wrst ": 134.5,
        "StDev ": 32.747
280     },

```

```

282         {
283             "count ": 11,
284             "host ": "lo0.fab4-1-gdc.ir2.yahoo.com",
285             "Loss%": 0.0,
286             "Snt ": 10,
287             "Last ": 130.62,
288             "Avg ": 46.746,
289             "Best ": 30.125,
290             "Wrst ": 130.62,
291             "StDev": 34.357
292         },
293         {
294             "count ": 12,
295             "host ": "usw1-1-lba.ir2.yahoo.com",
296             "Loss%": 0.0,
297             "Snt ": 10,
298             "Last ": 53.336,
299             "Avg ": 34.049,
300             "Best ": 30.023,
301             "Wrst ": 53.336,
302             "StDev": 8.066
303         },
304         {
305             "count ": 13,
306             "host ":
307                 "media-router-aol71.prod.media.vip.ir2.yahoo.com",
308             "Loss%": 0.0,
309             "Snt ": 10,
310             "Last ": 30.159,
311             "Avg ": 41.64,
312             "Best ": 30.008,
313             "Wrst ": 141.8,
314             "StDev": 35.2
315         }
316     ]
317 }

```

Wie zu erkennen ist wird durch diese z.B. die Hostnamen-Auflösungen übersprungen, was die Geschwindigkeit erhöht.

SS

netstat ist deprecated, es wird stattdessen dessen Nachfolger ss aus dem iproute2-Package verwendet:

```

Name       : iproute
2 Version   : 5.10.0
Release    : 2.fc34

```

```

4 Architecture : x86_64
  Size         : 1.7 M
6 Source       : iproute-5.10.0-2.fc34.src.rpm
  Repository   : @System
8 From repo    : anaconda
  Summary      : Advanced IP routing and network device
                  configuration tools
10 URL         : http://kernel.org/pub/linux/utils/net/iproute2/
  License      : GPLv2+ and Public Domain
12 Description : The iproute package contains networking
                  utilities (ip and rtmon,
                        : for example) which are designed to use the
                        : advanced networking
14             : capabilities of the Linux kernel.

```

Gehen Sie ins www und beobachten Sie die Veränderungen der netstat-Tabelle (netstat -an). Interpretieren Sie die Anzeige

Zuvor:

```

$ ss -tnp
2 State          Recv-Q          Send-Q          Local
                Address:Port      Address:Port      Process
FIN-WAIT-1      0                1
                10.60.54.18:60340
                104.17.239.204:443
4 FIN-WAIT-1      0                1
                10.60.54.18:52990
                104.16.18.94:443
ESTAB           0                0
                10.60.54.18:49524
                198.252.206.25:443
                users:(("chrome",pid=57314,fd=55))
6 FIN-WAIT-1      0                1
                10.60.54.18:48368
                151.101.1.69:443
FIN-WAIT-1      0                1
                10.60.54.18:45586
                142.250.186.161:443
8 FIN-WAIT-1      0                1
                10.60.54.18:60886
                151.101.14.217:443

```

```

FIN-WAIT-1          0          1

    10.60.54.18:45862
    23.185.0.3:443
10 ESTAB            0          0

    10.60.6.89:52008
    66.102.1.188:5228
    users:(( "chrome" ,pid=57314,fd=26))
FIN-WAIT-1          0          1

    10.60.54.18:42784
    104.244.42.193:443
12 FIN-WAIT-1       0          1

    10.60.54.18:43802
    140.82.121.3:443
FIN-WAIT-1          0          1

    10.60.54.18:56072
    104.19.154.83:443
14 ESTAB            0          0

    10.60.54.18:57766
    159.69.63.133:443
    users:(( "nextcloud" ,pid=4890,fd=38))
FIN-WAIT-1          0          1

    10.60.54.18:58314
    104.244.42.2:443
16 FIN-WAIT-1       0          1

    10.60.54.18:41736
    185.199.109.154:443

```

Nach dem Aufruf von news.ycombinator.com:

```

$ ss -tnp
2 State          Recv-Q      Send-Q      Local
                Address:Port  Process    Peer
                Address:Port
FIN-WAIT-1       0          1
                10.60.54.18:60340
                104.17.239.204:443
4 FIN-WAIT-1     0          1
                10.60.54.18:52990
                104.16.18.94:443

```

	ESTAB	0	0
	10.60.54.18:49524		
	198.252.206.25:443		
	users:(("chrome",pid=57314,fd=55))		
6	ESTAB	0	0
	10.60.6.89:50696		
	159.69.63.133:443		
	users:(("nextcloud",pid=4890,fd=65))		
	FIN-WAIT-1	0	1
	10.60.54.18:48368		
	151.101.1.69:443		
8	FIN-WAIT-1	0	1
	10.60.54.18:45586		
	142.250.186.161:443		
	FIN-WAIT-1	0	1
	10.60.54.18:60886		
	151.101.14.217:443		
10	FIN-WAIT-1	0	1
	10.60.54.18:45862		
	23.185.0.3:443		
	FIN-WAIT-2	0	0
	10.60.6.89:52008		
	66.102.1.188:5228		
12	FIN-WAIT-1	0	1
	10.60.54.18:56072		
	104.19.154.83:443		
	FIN-WAIT-1	0	1
	10.60.54.18:41736		
	185.199.109.154:443		
14	ESTAB	0	0
	10.60.6.89:50692		
	159.69.63.133:443		
	users:(("nextcloud",pid=4890,fd=38))		
	ESTAB	0	0
	10.60.6.89:47334		
	188.166.16.132:443		
	users:(("chrome",pid=57314,fd=40))		
16	FIN-WAIT-1	0	1

	10.60.54.18:54590 104.17.131.171:443 FIN-WAIT-1	0	1
18	10.60.54.18:53934 172.66.43.53:443 FIN-WAIT-1	0	1
	10.60.54.18:44820 185.199.111.133:443 FIN-WAIT-1	0	1
20	10.60.54.18:41740 185.199.109.154:443 ESTAB	0	0
	10.60.6.89:47336 188.166.16.132:443 users:(("chrome",pid=57314,fd=44)) FIN-WAIT-1	0	1
22	10.60.54.18:45360 104.17.211.204:443 ESTAB	0	0
	10.60.6.89:50686 159.69.63.133:443 users:(("nextcloud",pid=4890,fd=62)) FIN-WAIT-1	0	1
24	10.60.54.18:32944 151.101.13.132:443 ESTAB	0	0
	10.60.6.89:55356 209.216.230.240:443 users:(("chrome",pid=57314,fd=43)) FIN-WAIT-1	0	1
26	10.60.54.18:52794 66.102.1.188:5228 LAST-ACK	1	1
	10.60.54.18:37382 209.216.230.240:443 LAST-ACK	0	1043
28	10.60.54.18:57762 159.69.63.133:443 LAST-ACK	1	1

```

10.60.54.18:37378
209.216.230.240:443
FIN-WAIT-1      0      1

10.60.54.18:60308
151.101.12.193:443
30 ESTAB      0      0

10.60.6.89:50694
159.69.63.133:443
users:(("nextcloud",pid=4890,fd=63))
ESTAB      0      0

10.60.6.89:52010
66.102.1.188:5228
users:(("chrome",pid=57314,fd=26))
32 FIN-WAIT-1      0      1

10.60.54.18:41304
40.68.78.177:443
FIN-WAIT-1      0      1

10.60.54.18:38950
104.17.233.204:443
34 ESTAB      0      0

[2001:7c7:2121:8d00:1902:f308:6c8b:acb7]:50102
[2606:50c0:8001::153]:443
users:(("gnome-software",pid=4888,fd=92))
ESTAB      0      0

[2001:7c7:2121:8d00:1902:f308:6c8b:acb7]:50100
[2606:50c0:8001::153]:443
users:(("gnome-software",pid=4888,fd=42))

```

Wie zu sehen ist wurde eine TCP-Verbindung mit news.ycombinator.com aufgebaut:

```

$ dig +noall +answer news.ycombinator.com
2 news.ycombinator.com. 228 IN A 209.216.230.240

```

Testen Sie nun die Verbindung zwischen Ihrem PC und dem PC einer anderen Praktikumsgruppe und loten Sie die Möglichkeiten zur Verkehrsanalyse aus (netstat -s).

```

# Auf Host A
2 $ ss -tlnp
State      Recv-Q   Send-Q   Local Address:Port
Peer Address:Port   Process

```

```

4 LISTEN 0 128 0.0.0.0:22
  0.0.0.0:*
  LISTEN 0 1 0.0.0.0:6767
  0.0.0.0:* users:("nc",pid=10523,fd=3))
6 LISTEN 0 2 [::ffff:127.0.0.1]:3350
  *: *
  LISTEN 0 128 [::]:22
  [::]:*
8 LISTEN 0 2 *:3389
  *: *

$ nc -lp 6767
10 asdf

12 asdf
$ ss -tlnp
14 State Recv-Q Send-Q Local Address:Port Peer
    Address:Port Process
    LISTEN 0 128 0.0.0.0:22
    0.0.0.0:*
16 LISTEN 0 2 [::ffff:127.0.0.1]:3350
  *: *
  LISTEN 0 128 [::]:22
  [::]:*
18 LISTEN 0 2 *:3389
  *: *

20 # Auf Host B
$ ss -tnp | grep 6767
22 State Recv-Q Send-Q Local Address:Port Peer
    Address:Port Process
    ESTAB 0 0 141.62.66.5:54694
    141.62.66.4:6767 users:("nc",pid=36529,fd=3))
24 $ nc 141.62.66.4 6767
    asdf
26
    asdf
28 $ ss -tnp | grep 6767
    State Recv-Q Send-Q Local Address:Port
        Peer Address:Port Process

```

Wie zu Erkennen ist wurde eine TCP-Verbindung zwischen Host A und Host B erstellt, über welcher hier folgende Nachricht gesendet wurde:

```

1 asdf

3 asdf

```

Beobachten, dokumentieren und interpretieren Sie die Veränderungen der netstat-Tabelle beim „Durchklicken“ eines beliebigen Internet-Angebots.

```

1 $ ss -tnp
  State      Recv-Q      Send-Q          Local Address:Port
                        Peer  Address:Port    Process
3 $ ss -tnp
  State Recv-Q Send-Q Local Address:Port Peer
    Address:Port Process
5 ESTAB 0      0      141.62.66.5:54096      34.107.221.82:80
    users:(("firefox-esr",pid=36809,fd=98))
  ESTAB 0      0      141.62.66.5:52748      65.9.84.27:443
    users:(("firefox-esr",pid=36809,fd=41))
7 ESTAB 0      0      141.62.66.5:53806      54.239.39.102:443
    users:(("firefox-esr",pid=36809,fd=111))
  ESTAB 0      0      141.62.66.5:40840      142.250.186.138:443
    users:(("firefox-esr",pid=36809,fd=86))
9 ESTAB 0      0      141.62.66.5:36194      173.239.79.196:443
    users:(("firefox-esr",pid=36809,fd=77))
  ESTAB 0      0      141.62.66.5:33678      93.184.220.29:80
    users:(("firefox-esr",pid=36809,fd=34))
11 ESTAB 0      0      141.62.66.5:55186      162.219.226.52:443
    users:(("firefox-esr",pid=36809,fd=119))
  ESTAB 0      0      141.62.66.5:54384      209.216.230.240:80
    users:(("firefox-esr",pid=36809,fd=161))
13 ESTAB 0      0      141.62.66.5:36590      52.95.122.8:443
    users:(("firefox-esr",pid=36809,fd=141))
  ESTAB 0      0      141.62.66.5:46840      65.9.83.39:443
    users:(("firefox-esr",pid=36809,fd=74))
15 ESTAB 0      0      141.62.66.5:37550      54.239.39.102:80
    users:(("firefox-esr",pid=36809,fd=109))
  ESTAB 0      0      141.62.66.5:43074      142.250.185.67:80
    users:(("firefox-esr",pid=36809,fd=96))
17 ESTAB 0      0      141.62.66.5:54094      34.107.221.82:80
    users:(("firefox-esr",pid=36809,fd=85))
  ESTAB 0      0      141.62.66.5:42432      209.216.230.240:443
    users:(("firefox-esr",pid=36809,fd=172))
19 ESTAB 0      0      141.62.66.5:42430      209.216.230.240:443
    users:(("firefox-esr",pid=36809,fd=164))
  ESTAB 0      0      141.62.66.5:36288      65.9.83.11:443
    users:(("firefox-esr",pid=36809,fd=105))
21 ESTAB 0      0      141.62.66.5:50220      151.101.12.201:443
    users:(("firefox-esr",pid=36809,fd=84))
  ESTAB 0      0      141.62.66.5:42822      54.194.65.3:443
    users:(("firefox-esr",pid=36809,fd=120))
23 ESTAB 0      0      141.62.66.5:43710      2.21.21.24:80
    users:(("firefox-esr",pid=36809,fd=83))
  ESTAB 0      0      141.62.66.5:43922      54.68.102.210:443
    users:(("firefox-esr",pid=36809,fd=125))
25 ESTAB 0      0      141.62.66.5:42428      209.216.230.240:443
    users:(("firefox-esr",pid=36809,fd=162))
  ESTAB 0      0      141.62.66.5:42434      209.216.230.240:443
    users:(("firefox-esr",pid=36809,fd=176))

```

```

27 ESTAB 0      0      141.62.66.5:34436      162.219.224.163:443
      users:(("firefox-esr",pid=36809,fd=113))
ESTAB 0      0      141.62.66.5:44868      65.9.84.191:80
      users:(("firefox-esr",pid=36809,fd=140))
29 $ ss -tnp
State      Recv-Q      Send-Q      Local Address:Port
          Peer Address:Port      Process

```

Wie zu erkennen ist werden viele TCP-Verbindungen zu Webservern (Port 80 & Port 443) aufgebaut, hier zu news.ycombinator.com, eff.org und Amazon.

Route

route ist deprecated, es wird stattdessen ip route verwendet.

Interpretieren Sie die Einträge in der Routing-Tabelle Ihres Rechners.

Zu Erkennen ist dass das Default-Gateway 141.62.66.250 ist, über das Netzwerkgerät enp0s31f6. Auf localhost wird über den Kernel geroutet, d.h. dass Traffic niemals das System verlässt. Andere Subnetze werden über das Default-Gateway gerouted.

```

$ ip route show table all
2 default via 141.62.66.250 dev enp0s31f6
  141.62.66.0/24 dev enp0s31f6 proto kernel scope link src
    141.62.66.5
4 broadcast 127.0.0.0 dev lo table local proto kernel scope
  link src 127.0.0.1
  local 127.0.0.0/8 dev lo table local proto kernel scope host
  src 127.0.0.1
6 local 127.0.0.1 dev lo table local proto kernel scope host
  src 127.0.0.1
  broadcast 127.255.255.255 dev lo table local proto kernel
  scope link src 127.0.0.1
8 broadcast 141.62.66.0 dev enp0s31f6 table local proto kernel
  scope link src 141.62.66.5
  local 141.62.66.5 dev enp0s31f6 table local proto kernel
  scope host src 141.62.66.5
10 broadcast 141.62.66.255 dev enp0s31f6 table local proto
  kernel scope link src 141.62.66.5

```

Erweitern oder modifizieren Sie die Routing-Tabelle Ihres PC

Hier wurde nun eine neue Route hinzugefügt, welche das Subnet 192.0.2.128/25 über den Host 141.62.66.4 routed. Lädt der Host die richtigen Kernel-Module und wird IP-Forwarding mittels sysctl aktiviert, so könnte dieser damit als Router fungieren.

```

$ sudo ip route add 192.0.2.128/25 via 141.62.66.4
2 $ ip route show table all

```

```

    default via 141.62.66.250 dev enp0s31f6
4 141.62.66.0/24 dev enp0s31f6 proto kernel scope link src
    141.62.66.5
    192.0.2.128/25 via 141.62.66.4 dev enp0s31f6
6 broadcast 127.0.0.0 dev lo table local proto kernel scope
    link src 127.0.0.1
    local 127.0.0.0/8 dev lo table local proto kernel scope host
    src 127.0.0.1
8 local 127.0.0.1 dev lo table local proto kernel scope host
    src 127.0.0.1
    broadcast 127.255.255.255 dev lo table local proto kernel
    scope link src 127.0.0.1
10 broadcast 141.62.66.0 dev enp0s31f6 table local proto kernel
    scope link src 141.62.66.5
    local 141.62.66.5 dev enp0s31f6 table local proto kernel
    scope host src 141.62.66.5
12 broadcast 141.62.66.255 dev enp0s31f6 table local proto
    kernel scope link src 141.62.66.5

```

iperf

Mittels iperf3 kann die Übertragungsrate zwischen zwei Hosts getestet werden.

```

# Host A
2 $ iperf3 -s

```

```

4 Server listening on 5201

```

```

6 Accepted connection from 141.62.66.4, port 54336
[ 5] local 141.62.66.5 port 5201 connected to 141.62.66.4
    port 54338

```

ID	Interval	Transfer	Bitrate
[5]	0.00–1.00 sec	99.4 MBytes	834 Mb/s
[5]	1.00–2.00 sec	99.5 MBytes	835 Mb/s
[5]	2.00–3.00 sec	101 MBytes	846 Mb/s
[5]	3.00–4.00 sec	101 MBytes	845 Mb/s
[5]	4.00–5.00 sec	101 MBytes	845 Mb/s
[5]	5.00–6.00 sec	101 MBytes	844 Mb/s
[5]	6.00–7.00 sec	101 MBytes	844 Mb/s
[5]	7.00–8.00 sec	101 MBytes	850 Mb/s
[5]	8.00–9.00 sec	102 MBytes	853 Mb/s
[5]	9.00–10.00 sec	102 MBytes	856 Mb/s
[5]	10.00–10.00 sec	222 KBytes	756 Mb/s

```

20 [ ID] Interval          Transfer    Bitrate
22 # Host B
    $ sudo iperf3 -c 141.62.66.5
24 Connecting to host 141.62.66.5, port 5201
    [ 5] local 141.62.66.4 port 54338 connected to 141.62.66.5
    port 5201

```

26	[ID]	Interval		Transfer	Bitrate	Retr
			Cwnd				
	[5]	0.00–1.00	sec	101 MBytes	845 Mbites/sec	0
			342 KBytes				
28	[5]	1.00–2.00	sec	99.9 MBytes	838 Mbites/sec	0
			359 KBytes				
	[5]	2.00–3.00	sec	101 MBytes	845 Mbites/sec	0
			359 KBytes				
30	[5]	3.00–4.00	sec	101 MBytes	846 Mbites/sec	0
			359 KBytes				
	[5]	4.00–5.00	sec	101 MBytes	846 Mbites/sec	0
			359 KBytes				
32	[5]	5.00–6.00	sec	100 MBytes	840 Mbites/sec	0
			359 KBytes				
	[5]	6.00–7.00	sec	101 MBytes	844 Mbites/sec	0
			359 KBytes				
34	[5]	7.00–8.00	sec	101 MBytes	851 Mbites/sec	0
			359 KBytes				
	[5]	8.00–9.00	sec	102 MBytes	852 Mbites/sec	0
			359 KBytes				
36	[5]	9.00–10.00	sec	102 MBytes	859 Mbites/sec	0
			359 KBytes				

38	[ID]	Interval		Transfer	Bitrate	Retr
	[5]	0.00–10.00	sec	1009 MBytes	847 Mbites/sec	0
			sender				
40	[5]	0.00–10.00	sec	1008 MBytes	845 Mbites/sec	
			receiver				

42 iperf Done.

Hier kann z.B. erkannt werden, dass ca. 850 Mbites/sec erreicht werden können, was für die verwendete Gigabit-Netzwerkkarte mit CAT-5e-Kabel zu erwarten ist. ### NMAP Nmap ist die kurzform für Network Mapper. Mit diesem kann man Ports scannen, Informationen über die Services bekommen (Version, Betriebssystem etc.) und vorinstallierte als auch eigene Skripts verwenden.

Es gibt verschiedene Möglichkeiten Scans durchzuführen, der gängige (auch default) ist der TCP connect Port Scan. Es gibt noch weitere, welche situativ über Flags verwendet werden können:

```
$ nmap 10.10.247.15 -sS          # TCP SYN Port Scan
2 $ nmap 10.10.247.15 -sA        # TCP ACK Port Scan
$ nmap 10.10.247.15 -sU          # UDP Port Scan
```

Es besteht die Möglichkeit mehrere IPs zu scannen, ebenso wie ein Bereich von IPs, eine einzige IP oder eine Domain:

```
1 $ nmap 10.10.247.15            # Scannen einer
    einzigen IP
```

```

$ nmap 10.10.247.15 10.10.247.240      # Scannen mehrerer IPs
3 $ nmap 10.10.247.15-240              # Scannen des Bereichs
    von .15-.240
$ nmap scanme.nmap.org                # Scannen der Domain
    scanme.nmap.org

```

Es lassen sich ebenfalls die Ports definieren, welche auf einer IP gescannt werden sollen:

```

$ nmap 10.10.247.15 -p-                # Scannen der gesamten
    Portrange
2 $ nmap 10.10.247.15 -p 21            # Scannen des Port 21
$ nmap 10.10.247.15 -p 21-200         # Scannen alle Ports
    von 21 bis 200

```

Um Informationen bezüglich der verwendeten Versionen und Betriebssysteme zu erhalten können folgende Flags verwendet werden:

```

1 $ nmap 10.10.247.15 -sV              # Versucht die Version
    des Services zu ermitteln
$ nmap 10.10.247.15 -O                # Versucht das
    Betriebssystem zu ermitteln

```